

Introduction to the Sitar Framework

for Parallel Simulation of Synchronous Performance Models

Dr. Neha Karanjkar
Assistant Professor, CSE
Indian Institute of Technology Goa (IIT Goa)

nehak@iitgoa.ac.in

<https://nehakaranjkar.github.io>



Neha Karanjkar

Email: nehak@iitgoa.ac.in

Webpage:

<https://nehakaranjkar.github.io>



- **Education and Experience**

- M.Tech and Ph.D from IIT Bombay (EE)
- Research Scientist, AJIT processor development team, IITB
- Post-doc fellow at RBCCPS, IISc Bangalore
- Assistant Professor in CSE at IIT Goa since 2019

- **Research Areas:**

- Discrete-Event systems
- Parallel simulation
- Hybrid (Mixed Discrete-Continuous) Simulation
- Application Domains: Computer Systems Performance Modeling, Industrial Digital Twins and Process Modeling

Context, Overview and Purpose of this talk

- Types of models I'm going to focus on
- Need for Parallel Simulation and the reason that it can be Non-Trivial
- Sitar, its Features and Performance Results

The Need for Parallel Simulation

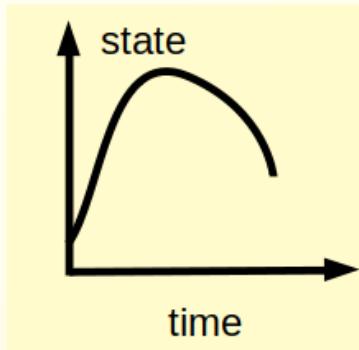
- **The Bad:** Models are getting larger, Benchmarks are getting longer
- **The Good:** Host systems are getting larger too! (many-cores, clusters, GPUs)
- **The Ugly:** There is something **fundamentally sequential** about dynamic simulation (in most cases). Parallelizing is non-trivial.

The Need for Parallel Simulation

- **The Obvious:** Design exploration can still be sped up (run many simulations in parallel). Single simulation time can still bottleneck design iterations!
- **The Opportunity:** Large models can be partitioned in domain-space for some parallelization. ...is linear scaling possible?
- **Technical Challenge:** Maintaining causality, deterministic execution.

Continuous System

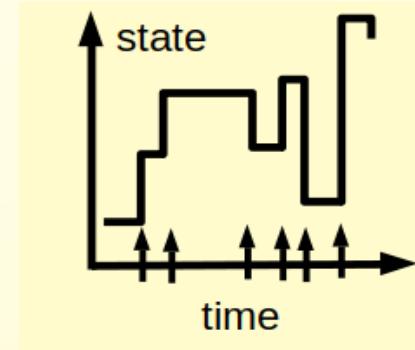
State of the system can evolve continuously with time



System behavior typically described as ODEs/PDEs

Discrete-Event System

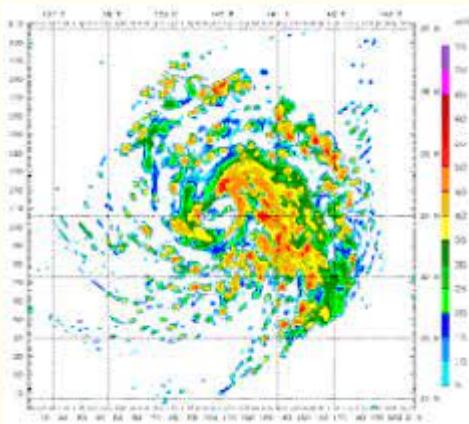
State changes assumed to happen only at discrete time-instants (called **Events**)



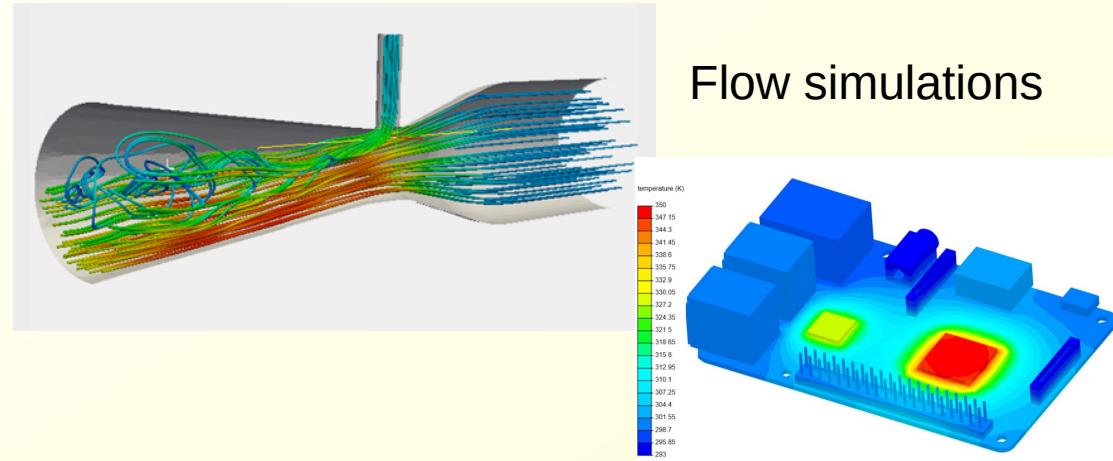
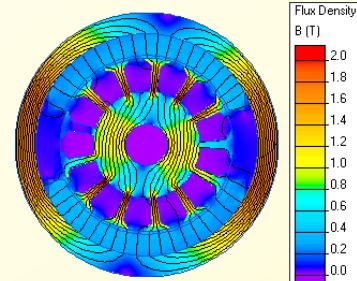
System behavior described as causality between events and time delays between them

Some Applications of Continuous System Simulations

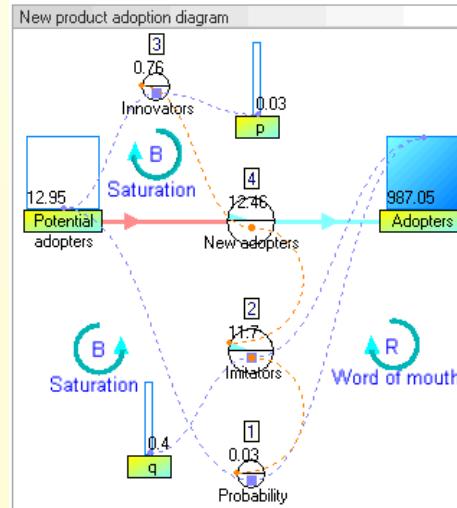
Weather simulation



Electro-mechanical simulation

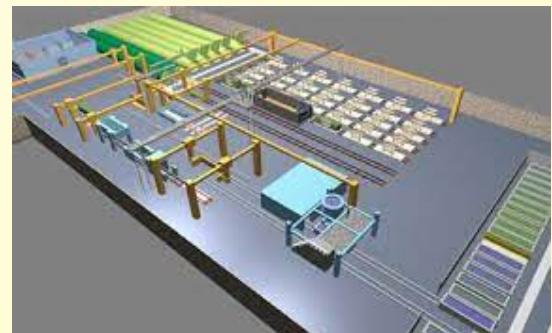
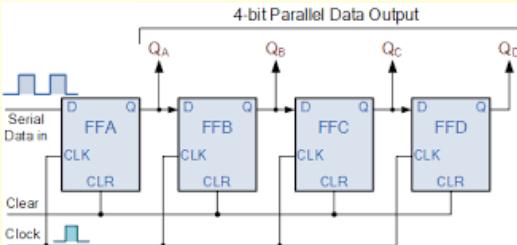
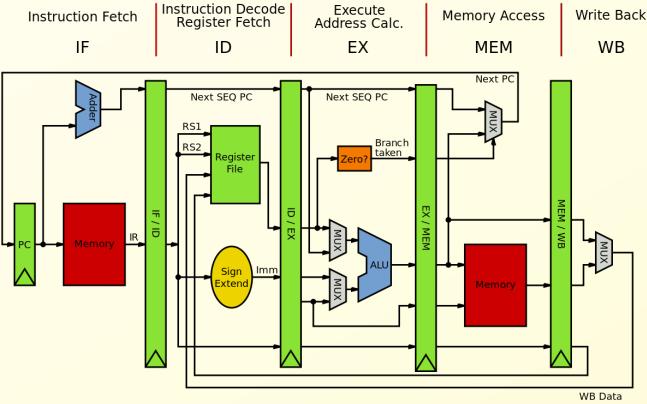
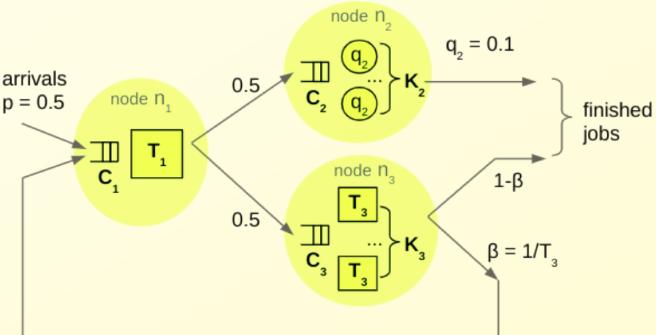
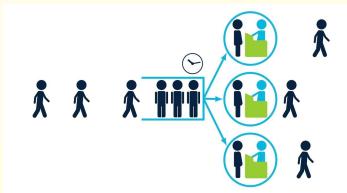
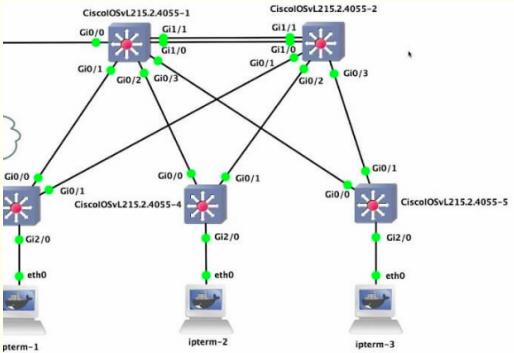


System Dynamics Simulation



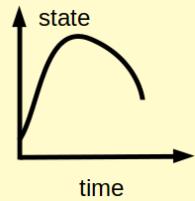
Discrete-Event Simulation Applications

- Computer networks
- Digital systems
- Computer architecture
- Queueing networks
- Factory/Assembly lines
- Models of customer arrival, behavior, inventory, sales etc in retail, airports, restaurants ..



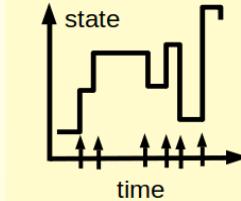
Continuous System Model

State of the system is assumed to evolve **continuously** over time



Discrete-Event System Model

State changes assumed to happen only at **discrete time-instants** (called **Events**)



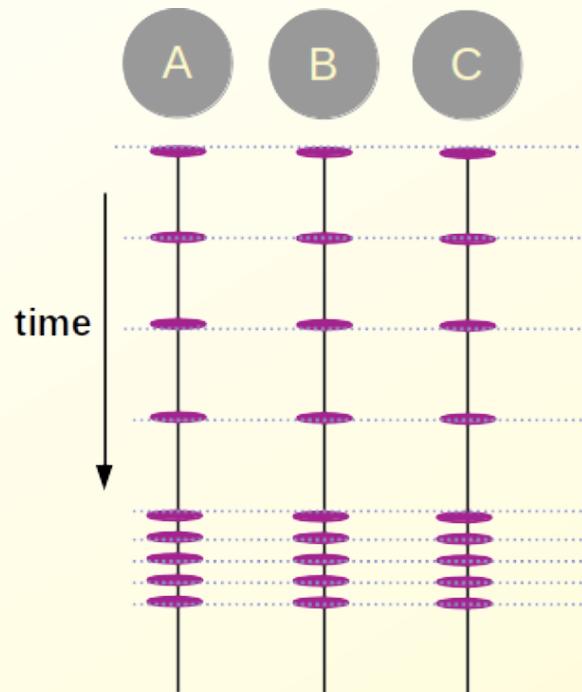
The difference is conceptual. In practice...

- Continuous simulators (on digital computers) **have to** discretize the state-space and time axis.
- The body of techniques and approaches, challenges, history, and domain experts are generally separate between these two areas.

Continuous Simulation

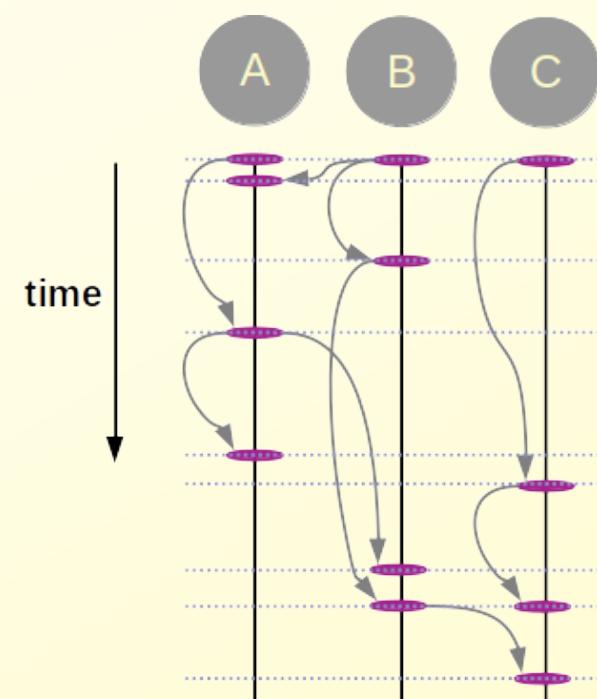
discretization

Time-stepped Approach



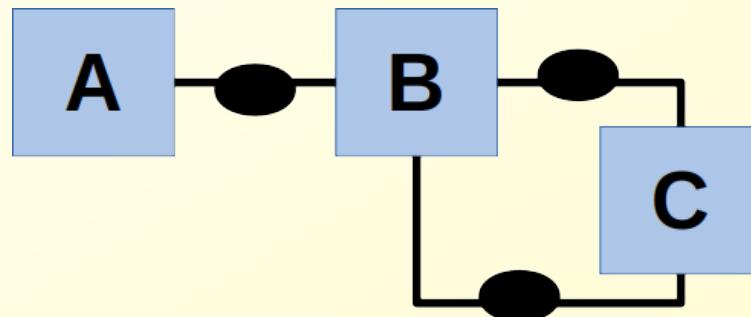
Discrete-Event Simulation

Event-driven Approach



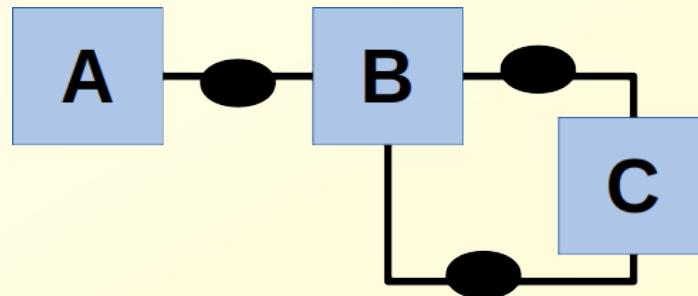
Synchronous / Cycle-based Models

- All components operate on a single global “clock”. All event times are integer-valued
- System can consist of concurrent components with static or dynamic communication structures



Cycle-based Simulation

- Execute the behavior of each component in every clock cycle?
- Order of execution among modules matters!
- Race conditions on updating communication channels?

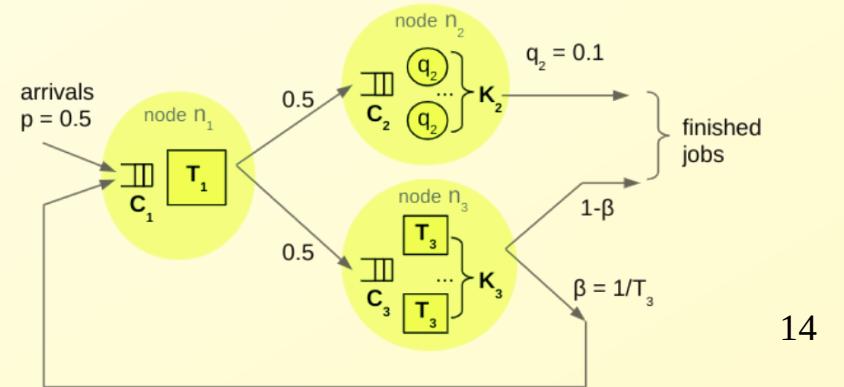
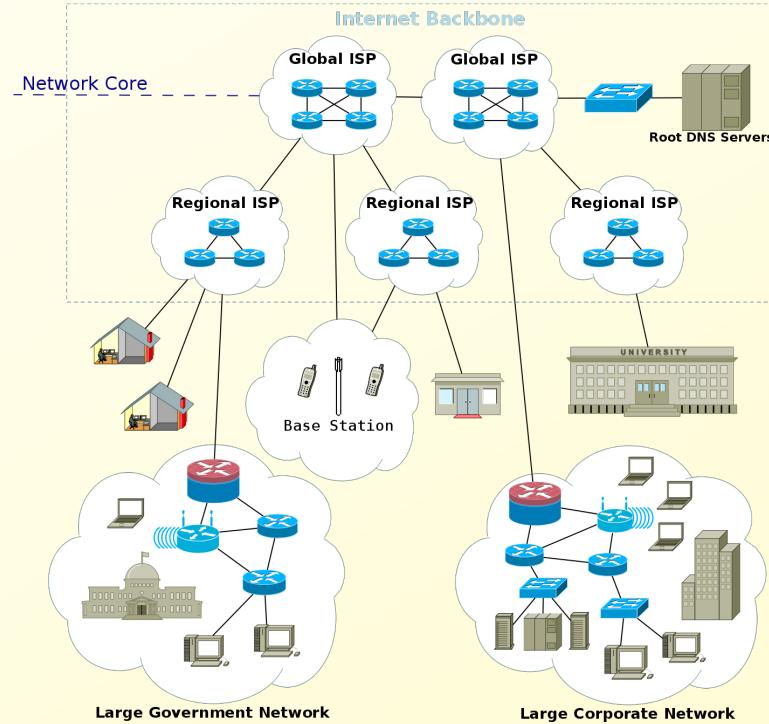
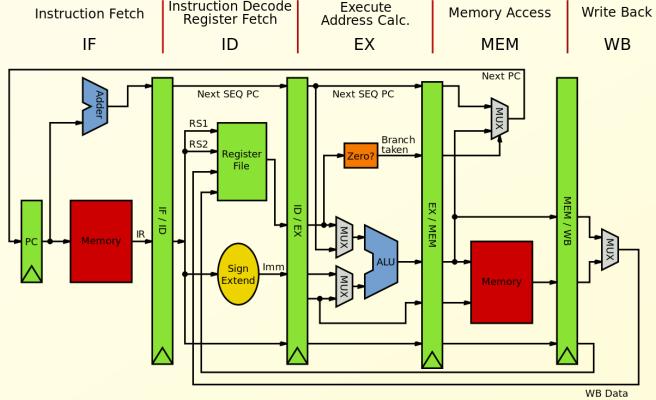
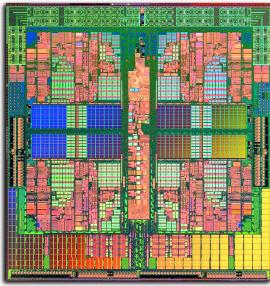


Sitar: A modeling and parallel simulation framework targeted for systems ...

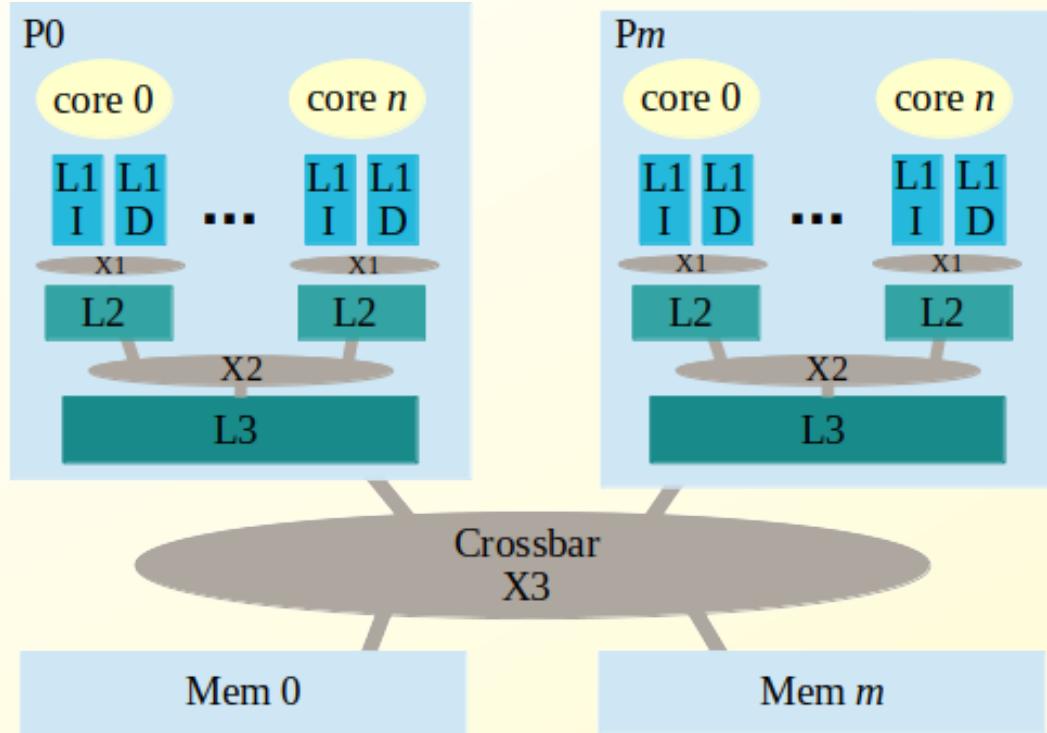
- that are large, with many concurrent components having complex, distinct behaviors,
- have a static structure (hierarchy and interconnections),
- operate in discrete-time, and
- where the communication between components incurs some delay...

Some Examples:

- Computer Networks
- Architecture-level models for processors and Systems-on-Chip
- Clocked VLSI systems
- Discrete-time Queueing Networks
- ...



Images sourced from Wikimedia commons



About Sitar

- An open-source cycle-based modeling framework
- Targeted for writing quick architecture-level models for fast, parallel simulation on shared-memory systems
- Consists of a Modeling Language and a Cycle-based Simulation kernel

Development

- Version 1 Developed as an internal tool for computer architecture research at IIT Bombay. (2013)
- Used in multi-core design exploration studies and simulating and optimizing queueing networks (2013-2015)
- Version 2: Open-source (MIT-License)

<https://sitar-sim.github.io/sitar/>

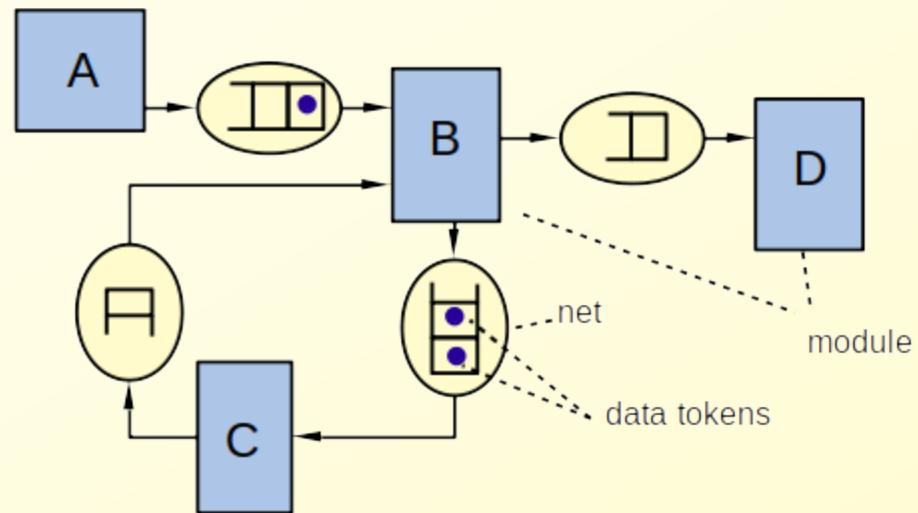
The screenshot shows a screenshot of a web browser displaying the GitHub page for the Sitar simulation framework. The URL in the address bar is nehakaranjkar.github.io/sitar/. The page title is "Sitar" and it is described as "A cycle-based simulation framework". There are three download buttons: "Download .zip", "Download .tar.gz", and "View on GitHub". Below the download buttons is a brief description of the Sitar framework, mentioning its purpose as a modeling and simulation tool for cycle-based systems, its language, and some of its features. At the bottom, there is a list of constructs supported by the language.

Sitar is a framework for modeling and simulation of cycle-based (discrete-time) systems. It consists of a system description language and a lightweight simulation kernel. A system is described as a set of concurrent, interconnected modules operating on a global clock. The language supports hierarchical descriptions: modules can contain instances of other modules. The behavior of each module can be described in an imperative manner as a sequence of statements. The language supports constructs such as :

- Time delays
- Parallel blocks
- Branch/loop constructs
- Procedures and
- Code blocks (instantaneous)

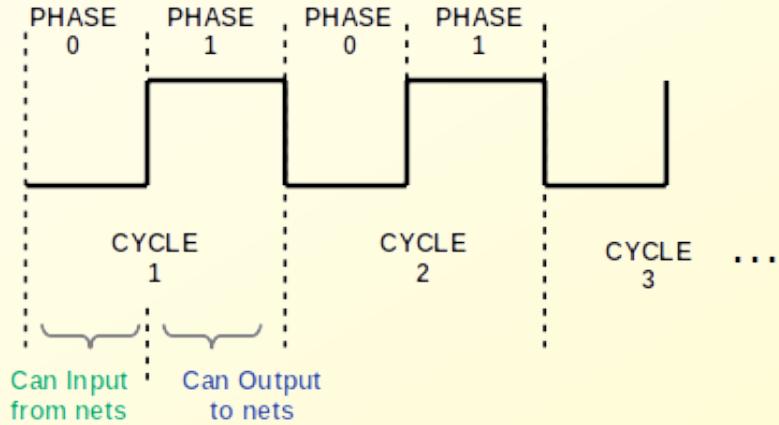
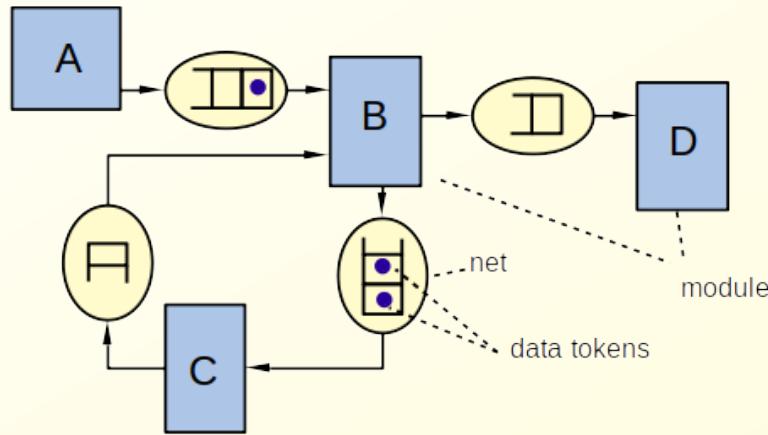
Sitar Execution Model

- A system consists of *modules* communicating over FIFO channels (*nets*) by transfer of *data-tokens*
- All modules operate on a single clock



Sitar Execution Model

- Two-phase execution (deterministic) : input in phase 0, output in phase 1



- **Restriction:** Delay of at-least one cycle for information to propagate from the input of a module to it's output. (Communicating Moore machines paradigm)

Sitar Execution Model

- Two-phase simulation algorithm (pseudo-code)

```
cycle = 0
while (cycle < simulation_end_time)
{
    phase=0;//run all modules in phase 0
    for (m=0; m<num_modules; m++)
        module[m].run(cycle, phase);

    phase=1;//run all modules in phase 1
    for (m=0; m<num_modules; m++)
        module[m].run(cycle, phase);

    cycle = cycle + 1
}
```

Sitar Execution Model

- Two-phase simulation algorithm - Parallelized using OpenMP

```
cycle = 0
while (cycle < simulation_end_time)
{
    phase=0;
    #pragma omp for //run all modules for phase 0
    for (m=0; m<num_modules; m++)
        module[m].run(cycle, phase);
    #pragma omp barrier

    phase=1;
    #pragma omp for //run all modules for phase 1
    for (m=0; m<num_modules; m++)
        module[m].run(cycle, phase);
    #pragma omp barrier
    cycle = cycle + 1
}
```

Sitar Modeling Language

- **Describing Structure:** Interconnection of modules and nets, Hierarchy, Generics (templated modules), For-loops to create regular structures
- **Describing Module Behavior:** In an imperative manner with conditional wait statements and delays, branches, loops, fork-join parallel blocks and embedded code
- **Other features:** Logging support, debugging support, syntax highlighting
- **Translation:** Each module description gets translated to a C++ class, compiled together with a lightweight simulation kernel to get an executable.

Sitar Modeling Language: Structure

```
module Top
    //Instantiate a shift register
    //with 3 stages, and a per-stage delay of 1
    submodule S : ShiftRegister<3,1>

    //Instantiate the Producer and Consumer
    submodule P : Producer
    submodule C : Consumer

    //Connect the Producer and Consumer
    //to the two ends of the ShiftRegister
    P.out_port => S.n[0]
    C.in_port  <= S.n[3]
end module
```

```
module Stage
    parameter int DELAY = 1
    import in_port
    export out_port
behavior
    // describe the behavior of each stage
    //...
end behavior
end module
```

```
module ShiftRegister
    //declare module parameters
    //and their default values
    parameter int N = 1      //number of stages
    parameter int DELAY = 1   //delay of each stage

    //The ShiftRegister consists of N stages
    submodule_array stage[N] : Stage<DELAY>
    //N+1 nets to connect all stages
    net_array n[N+1] : capacity 1

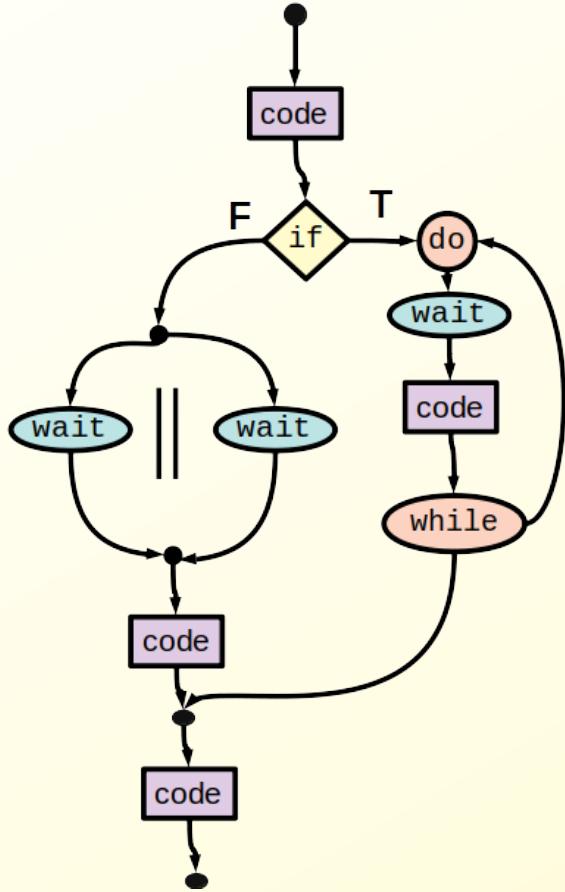
    //connect the stages via nets
    for i in 0 to (N - 1)
        stage[i].in_port  <= n[i]
        stage[i].out_port => n[i+1]
    end for
end module
```

Sitar Modeling Language: Behavior

Module Behavior: described in an imperative manner as a sequence of statements

- **Atomic statements (instantaneous)**
 - wait (time), wait until (condition)
 - C++ code blocks
 - simulation control, logging
- **Compound statements (in-turn contain a sequence)**
 - do-while, if-else
 - Parallel blocks (fork-join concurrency)
 - Procedures

Modeling Language: Behavior



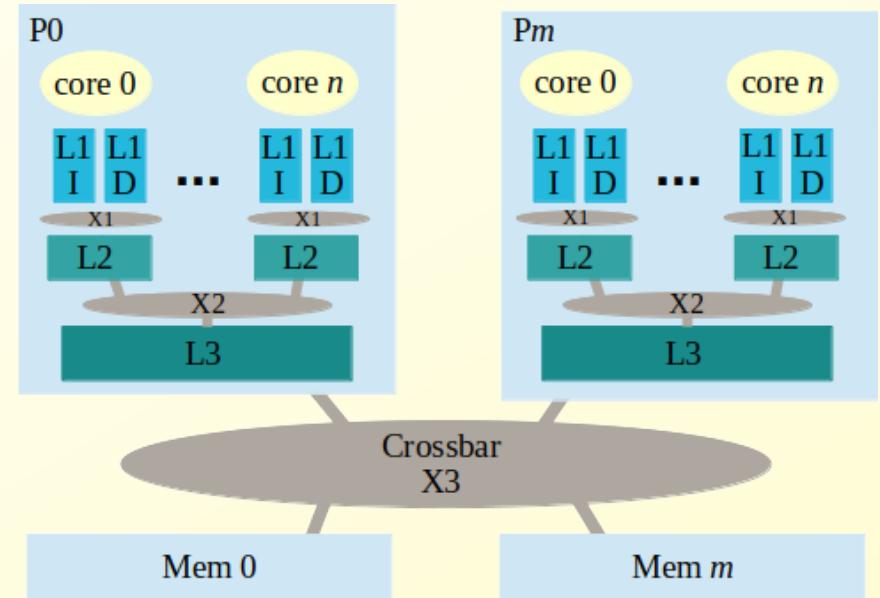
```
module Top
behavior

    decl $int x;$;
    //Input x from user
    $ std::cout<<"\nEnter a number:";std::cin>>x;$;

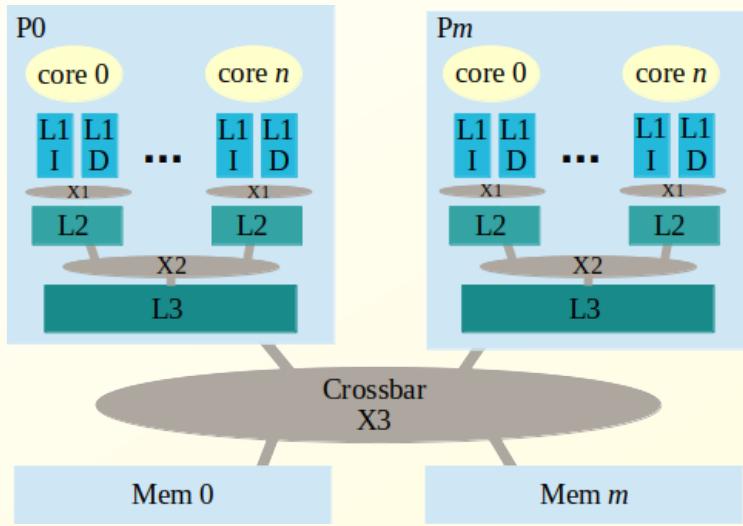
    //Branch
    if($x%2==0$) then
        do //Loop
            wait(1,0); //wait for one cycle
            $log<<endl<<"In the even branch";$;
            while($current_time.cycle()<4$) end do;
            stop simulation;
    else
        [ //Parallel block
            wait (2,0);
        ||
            wait until ($current_time>=time(x,0$));
        ];
        $log<<endl<<"In the odd branch";$;
    end if;
    $log<<endl<<"Done";$;
    stop simulation;
end behavior
end module
```

Use Case and Scalability

- Parameterized Multi-core model with 3-level cache hierarchy
- Cores: in-order, Sparc V8 ISA
- Caches: per-core split L1, unified L2, shared L3
- Directory-based hierarchical MESI protocol
- Can run Sparc v8 binaries



Use Case and Scalability

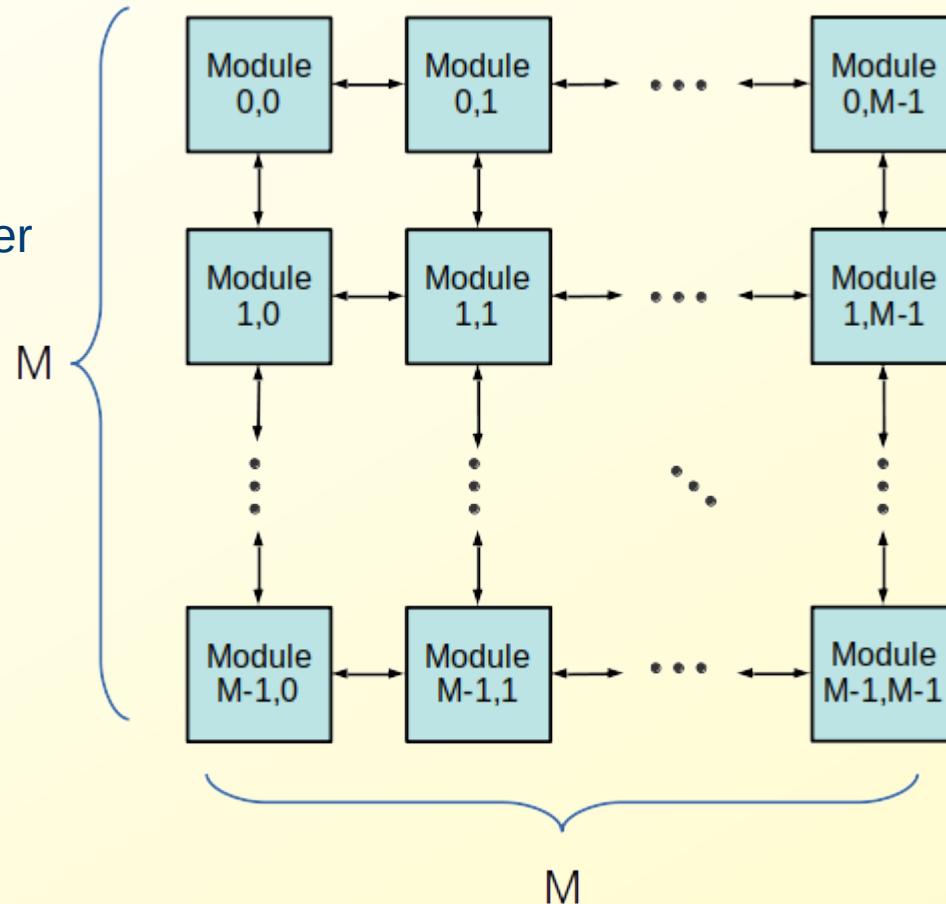


Kernel	Problem Size
Embarrassingly Parallel (EP)	2^{16}
Multigrid (MG)	16^3
3-D FFT PDE solver (FT)	32^3
Integer Sort (IS)	$2^{14} + 2^{12}$

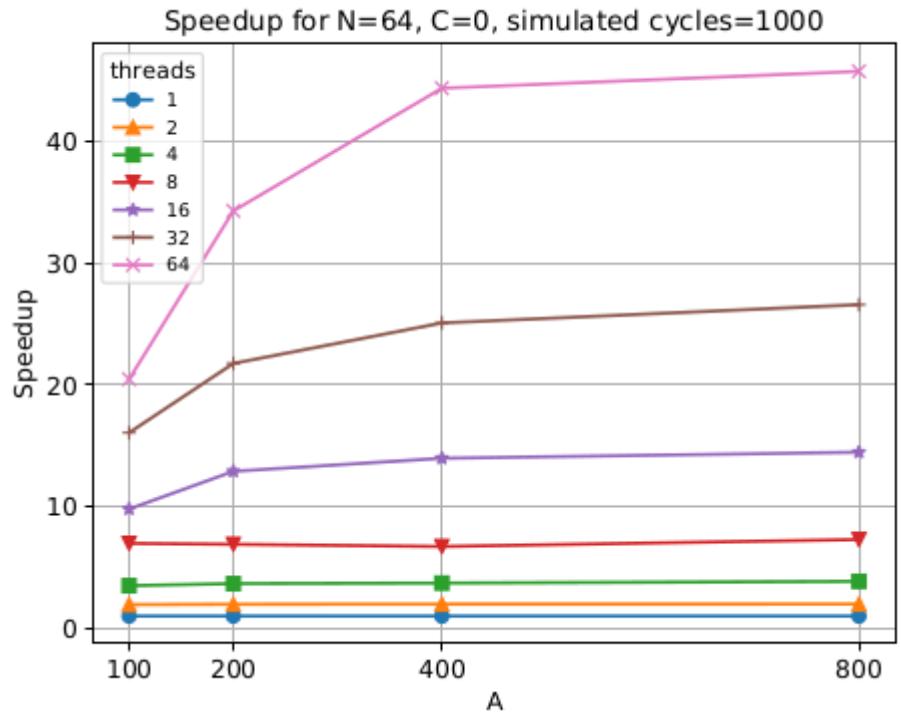
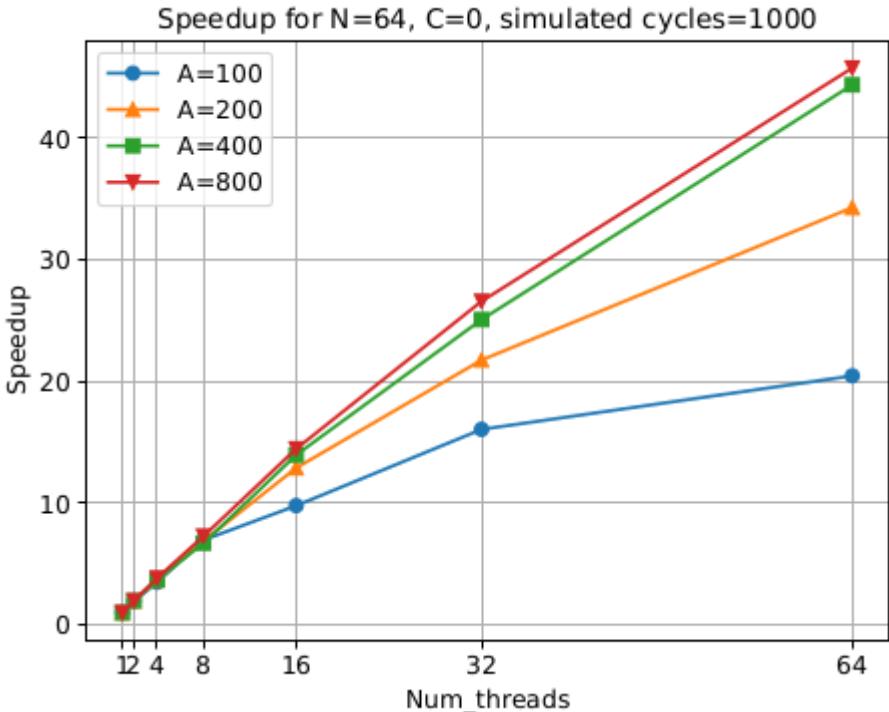
Workload (Sim Cycles)	Simulation time in seconds and (Speedup)		
	1 thread	2 threads	4 threads
EP(9781300)	954 (1x)	622 (1.5x)	442 (2.2x)
FT(16578568)	1710 (1x)	1112 (1.5x)	855 (2x)
IS(1434771)	156 (1x)	102 (1.5x)	78 (2x)
MG(13728356)	1452 (1x)	944 (1.5x)	726 (2x)

Parameterized Benchmark Model

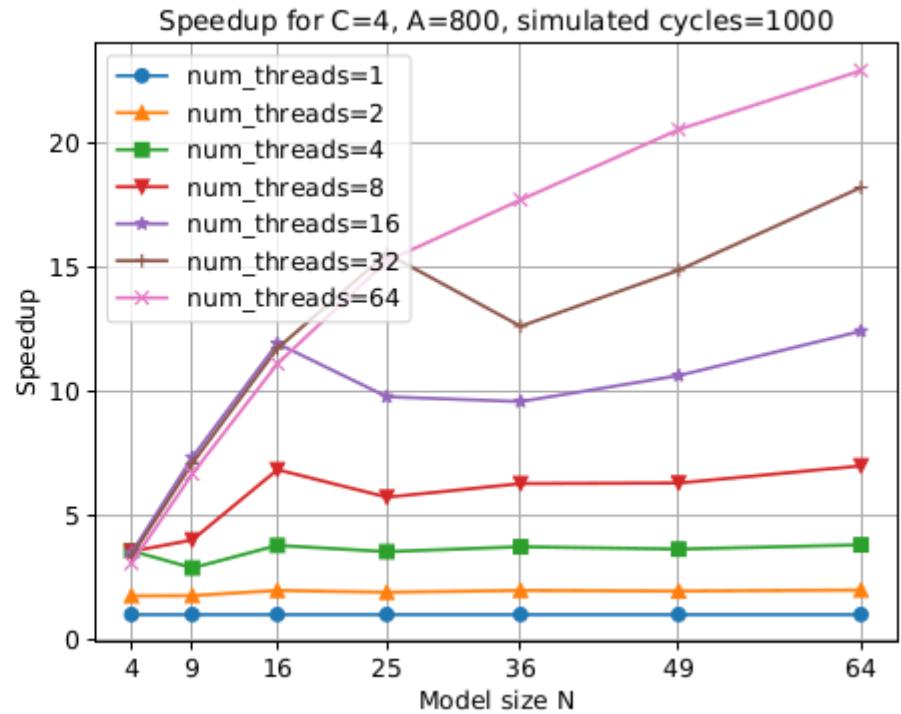
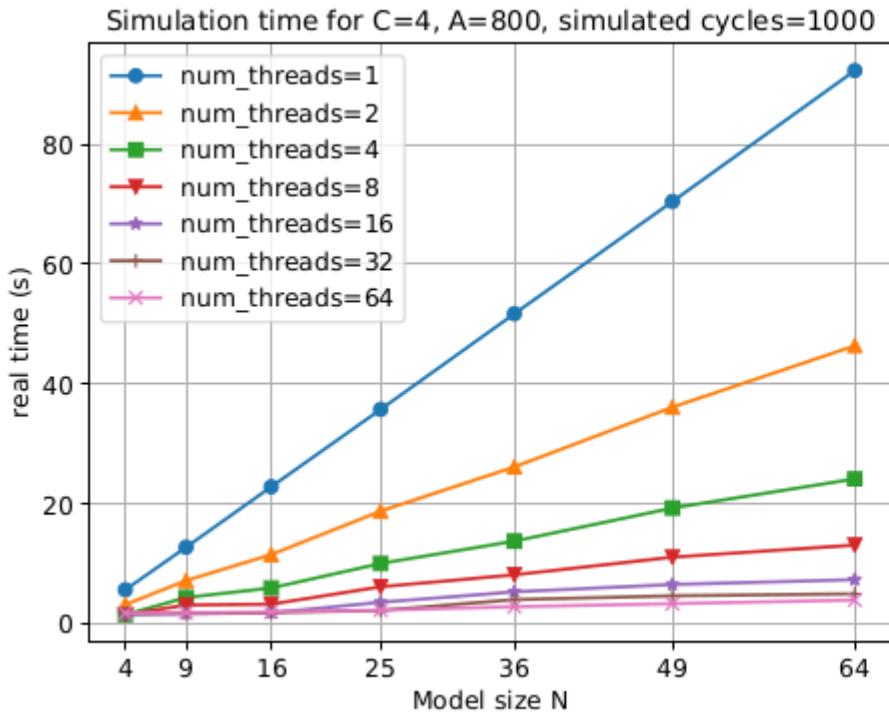
- $N = M \times M$: Number of modules in the system
- A : Computation parameter
- C : Communication parameter



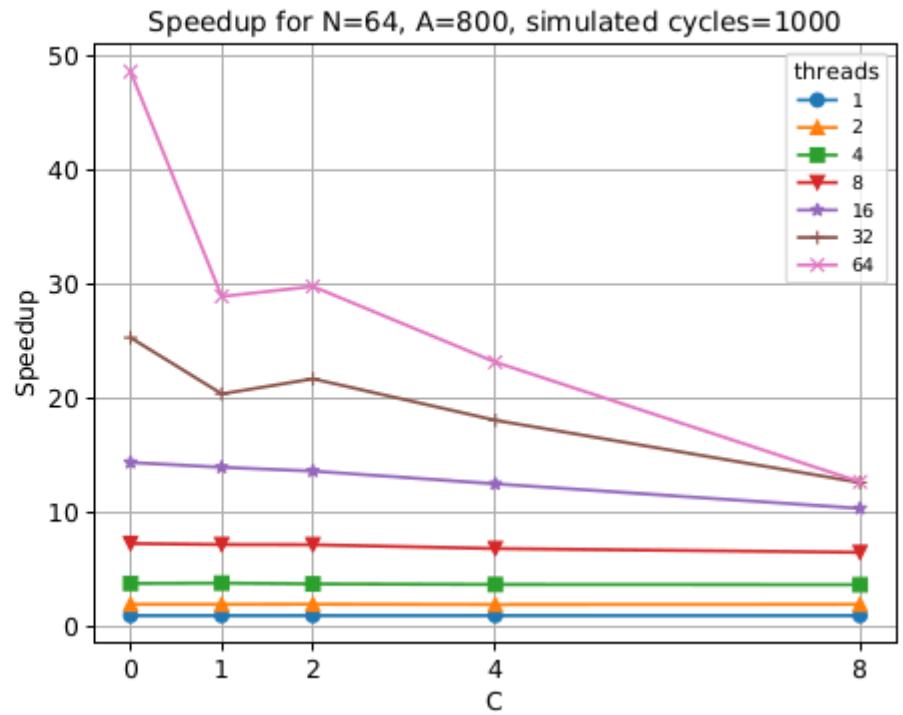
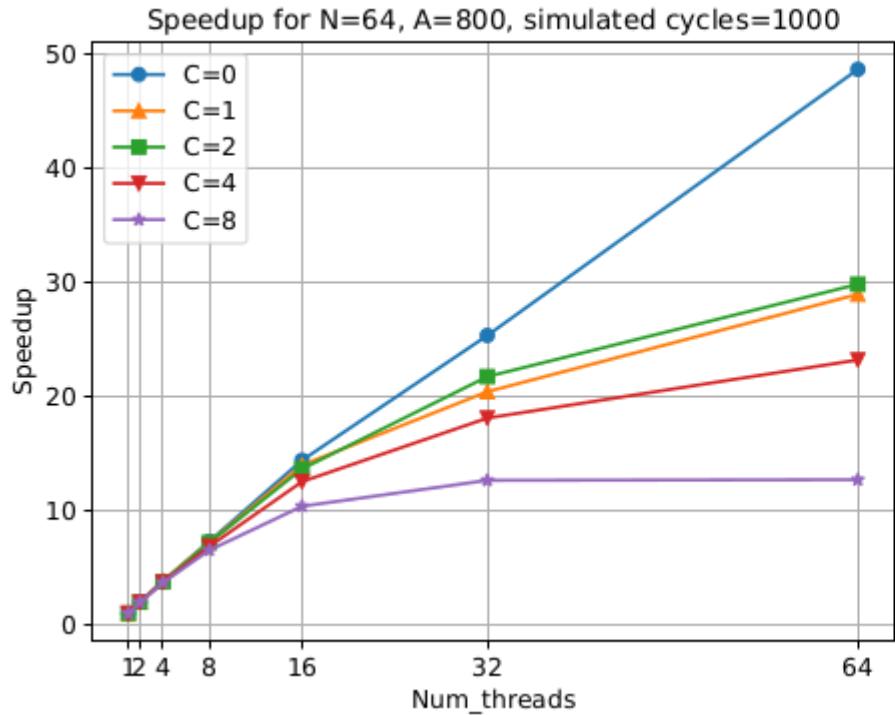
Scalability (Model: Mesh with N modules Host:40-core rack server)



Scalability (Model: Mesh with N modules Host:40-core rack server)

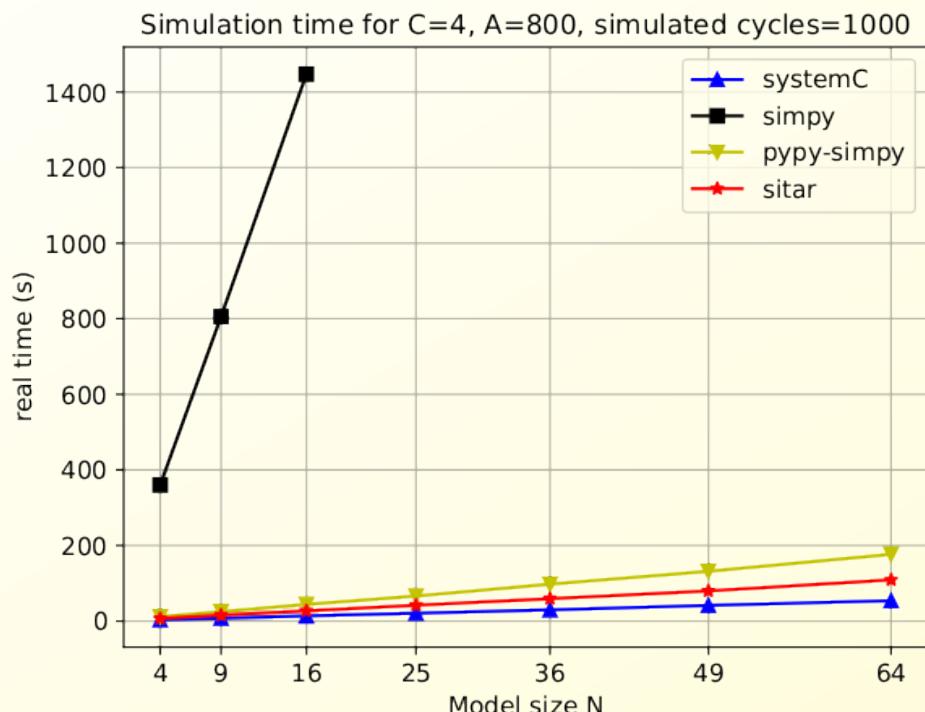


Scalability (Model: Mesh with N modules Host:40-core rack server)

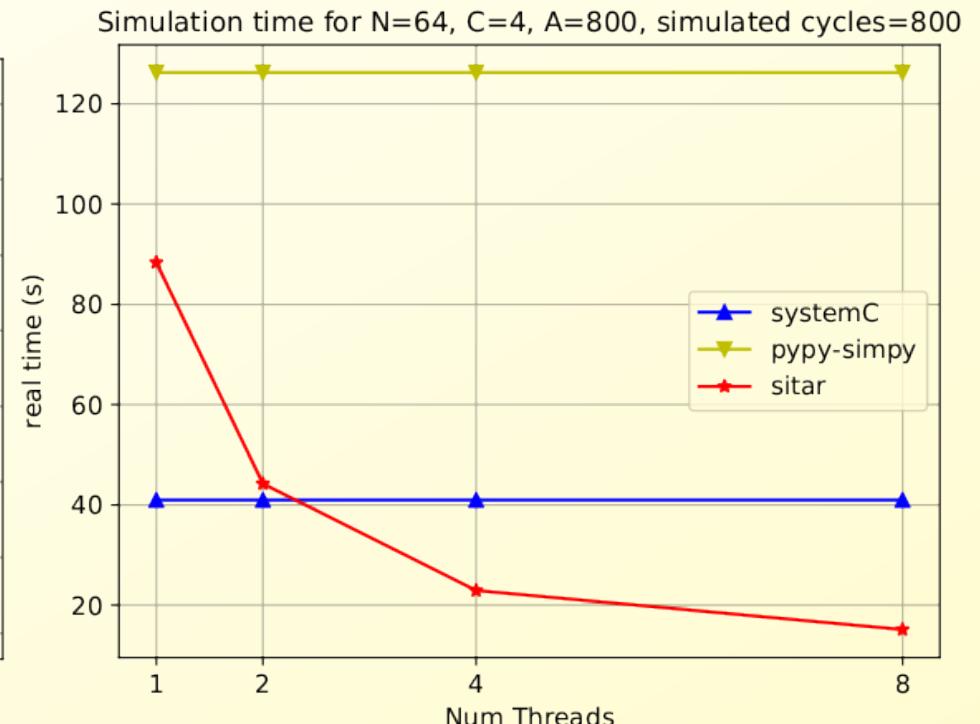


Performance Comparison with SimPy and SystemC

Single-threaded



Multi-threaded Sitar



Ongoing work

- Shared variable access to nets - performance takes a hit due to cache line transfer between threads.
- To be explored: MPI based communication and distributed system simulation
- Python port

Conclusions

- Cycle-based simulation algorithm of Sitar enables fast, parallel simulations for discrete-time systems
- Sitar's modeling language provides a rich set of features for describing structure and behavior, along with logging support
- Sitar is suited to modeling discrete-time systems with a structural aspect (such as computer architecture models, networks etc)

References

- [1] **Sitar: A Cycle-based Discrete-Event Simulation Framework for Architecture Exploration [Best Paper Award]** Neha Karanjkar, Madhav Desai, *Proceedings of the 12th International Conference on Simulation and Modeling Methodologies, Technologies and Applications (SIMULTECH 2022)*, July 2022, Lisbon, Portugal
- [2] **Efficient Parallel Simulation of Networked Synchronous Discrete-Event Systems**, Neha Karanjkar, Madhav Desai, Akhil Kushe, and Anish Natekar, Accepted for publication at WinterSim 2024, Orlando, Florida USA Dec 2024